



UNIVERSIDAD NACIONAL DE SAN AGUSTIN

ESCUELA PROFESIONAL DE  
CIENCIA DE LA COMPUTACIÓN  
ALGORITMOS PARALELOS

---

CUDA

---

***Alumna :***

*Rosa Yuliana Gabriela*

*Paccotacya Yanque*

***Profesor:***

*Mg. Alvaro Henry Mamani*

*Aliaga*

## Índice

<b>1. Suma de Vectores</b>	<b>2</b>
1.1. Código . . . . .	2
1.2. Output . . . . .	3
<b>2. Suma de Matrices</b>	<b>3</b>
2.1. Código . . . . .	3
2.2. Output . . . . .	5
<b>3. Procesamiento de Imágenes</b>	<b>5</b>
3.1. Escala de Grises . . . . .	5
3.2. Blur . . . . .	7

Todas las pruebas y ejecuciones de código se realizaron en una PC con NVIDIA Corporation GK208M GeForce GT 740M con capability de 3.0 y Cuda 8.0.

## 1. Suma de Vectores

### 1.1. Código

```
1 #include <iostream>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <vector>
5 #include <string>
6 #include <cuda.h>
7 using namespace std;
8 __global__ void suma_vectores(float *c ,float *a , float *b,int N)
9 {
10     int idx=blockIdx.x * blockDim.x+ threadIdx.x;
11     if(idx<N)
12     {
13         c[idx]=a[idx] + b[idx];
14     }
15 }
16
17 int main(void)
18 {
19     float *a_h,*b_h,*c_h;
20     float *a_d,*b_d,*c_d;
21     int N=1000000;
22     size_t size=N*sizeof(float);
23     a_h = (float *) malloc (size);
24     b_h = (float *) malloc (size);
25     c_h = (float *) malloc (size);
26     srand(1);
27
28     for (int i = 0; i < N; ++i) {
29         a_h[i] = rand();
30         b_h[i] = rand();
31     }
32
33
34     cudaMalloc((void**)& a_d, size);
35     cudaMalloc((void**)& b_d, size);
36     cudaMalloc((void**)& c_d, size);
37
38     cudaMemcpy(a_d, a_h, size ,cudaMemcpyHostToDevice);
```

```

39  cudaMemcpy(b_d, a_h, size, cudaMemcpyHostToDevice);
40  int block_size=256;
41  int n_blocks=N/block_size + (N%block_size ==0 ? 0:1);
42  suma_vectores <<< n_blocks, block_size >>> (c_d, a_d, b_d, N);
43  cudaMemcpy(c_h, c_d, size, cudaMemcpyDeviceToHost);
44  printf("Suma con  %d hebras con  %d hebras por bloque!\n", N,
        block_size);
45
46  /*for (int i=0;i<N;i++)
47  {
48      cout<<c_h[i]<<" "<<endl;
49  }*/
50
51  free(a_h);
52  free(b_h);
53  free(c_h);
54  return(0);
55 }

```

Código 1: Suma de vectores

## 1.2. Output

```

rose@Satellite-S55-A:~/CS_AlgoritmosParalelos/cuda$ nvcc -o a sum_vect.cu
nvcc warning : The 'compute_20', 'sm_20', and 'sm_21' architectures are deprecated, and may be removed in a future release (Use -Wno-deprecated-gpu-targets to suppress warning).
rose@Satellite-S55-A:~/CS_AlgoritmosParalelos/cuda$ ./a
Suma con  1000000 hebras con 256 hebras por bloque!

```

## 2. Suma de Matrices

### 2.1. Código

```

1  #include <cuda.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <stdint.h>
5
6  const int N = 16384;
7  const int THREADS_PER_BLOCK = 512;
8
9  __global__ void add_threads_blocks (int *a, int *b, int *c, int n) {
10
11  int index = threadIdx.x * blockDim.x * threadIdx.x;

```

```
12     if (index < n) {
13         c[index] = a[index] + b[index];
14     }
15 }
16
17 int main(void) {
18     int *a, *b, *c;
19     int *d_a, *d_b, *d_c;
20     size_t size = N * sizeof(int);
21
22     srand(1);
23
24     a = (int *) malloc(size);
25     b = (int *) malloc(size);
26     c = (int *) malloc(size);
27
28
29     for (int i = 0; i < N; ++i) {
30         a[i] = rand();
31         b[i] = rand();
32     }
33     //uint kernelTime;
34     //cutCreateTimer(&kernelTime);
35     //cutResetTimer(kernelTime);
36     cudaMalloc((void **) &d_a, size);
37     cudaMalloc((void **) &d_b, size);
38     cudaMalloc((void **) &d_c, size);
39     cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
40     cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
41     //cutStartTimer(kernelTime);
42     add_threads_blocks<<<(N + (THREADS.PER.BLOCK - 1)) /
43         THREADS.PER.BLOCK, THREADS.PER.BLOCK>>>(d_a, d_b, d_c, N);
44     // cudaThreadSynchronize();
45     //cutStopTimer(kernelTime);
46     cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);
47     printf("Suma con %d hebras con %d hebras por bloque!\n", N,
48         THREADS.PER.BLOCK);
49     //printf("Time for the kernel: %f ms\n", cutGetTimerValue(kernelTime));
50     free(a); free(b); free(c);
51     cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
52     return 0;
53 }
```

Código 2: Suma de matrices

## 2.2. Output

```
rose@Satellite-S55-A:~/CS_AlgoritmosParalelos/cuda$ nvcc -o a sum_matrix.cu
nvcc warning : The 'compute_20', 'sm_20', and 'sm_21' architectures are deprecated, and may be removed in a future release (Use -Wno-deprecated-gpu-targets to suppress warning).
rose@Satellite-S55-A:~/CS_AlgoritmosParalelos/cuda$ ./a
Suma con 16384 hebras con 512 hebras por bloque!
```

## 3. Procesamiento de Imágenes

Para obtener las matrices de las imagenes y tambien poder mostrarlas luego de su procesamiento se usó Cimg y archivos .dat.

### 3.1. Escala de Grises

```
1 #include <stdio.h>
2 #include <fstream>
3 #include <iostream>
4 #define CHANNELS 3 // we have 3 channels corresponding to RGB
5 using namespace std;
6
7 #define CHANNELS 3 // we have 3 channels corresponding to RGB
8 // The input image is encoded as unsigned characters [0, 255]
9 --global-- void colorConvert(float * Pout, float * Pin, int width, int
    height)
10 {
11     int Col = threadIdx.x + blockIdx.x * blockDim.x;
12     int Row = threadIdx.y + blockIdx.y * blockDim.y;
13     if (Col < width && Row < height)
14     {
15         // get 1D coordinate for the grayscale image
16         int greyOffset = Row*width + Col;
17         // one can think of the RGB image having
18         // CHANNEL times columns than the grayscale image
19         int rgbOffset = greyOffset*CHANNELS;
20         float r = Pin[rgbOffset]; // red value for pixel
21         float g = Pin[rgbOffset + 1]; // green value for pixel
22         float b = Pin[rgbOffset + 2]; // blue value for pixel
23         // perform the rescaling and store it
24         // We multiply by floating point constants
25         Pout[greyOffset] = 0.21f*r + 0.71f*g + 0.07f*b;
26     }
```

```

27 }
28
29
30 void save_data(float o[225][225])
31 {
32     ofstream archivo("gray.dat");
33     for (int i = 0; i < 225; ++i)
34     {
35         for (int j = 0; j < 225; ++j)
36         {
37             archivo<<o[i][j]<<" ";
38         }
39         archivo<<endl;
40     }
41 }
42
43 void GrayScale(float m[225][225*3], int width, int height)
44 {
45     float o[225][225];
46
47     int size_in = width * (height*3);
48     int size_out = width * height;
49     int memSize_in = size_in * sizeof(float);
50     int memSize_out = size_out * sizeof(float);
51
52     float *d_A, *d_B;
53
54     cudaMalloc((void **) &d_A, memSize_in);
55     cudaMalloc((void **) &d_B, memSize_out);
56
57     cudaMemcpy(d_A, m, memSize_in, cudaMemcpyHostToDevice);
58
59     dim3 DimGrid(floor((width-1)/16 + 1), floor((height-1)/16+1), 1);
60     dim3 DimBlock(16, 16, 1);
61     colorConvert<<<DimGrid,DimBlock>>>(d_B, d_A, width, height);
62
63     cudaMemcpy(o, d_B, memSize_out, cudaMemcpyDeviceToHost);
64
65     cudaFree(d_A);
66     cudaFree(d_B);
67     save_data(o);
68 }
69
70 void leer_data(const char *file, float m[225][225*3])
71 {
72     char buffer[100];
73     ifstream archivo2("image.dat");
74     for (int ii = 0; ii < 225; ++ii)
75     {

```

```
76     for (int jj = 0; jj < 225; ++jj)
77     {
78         archivo2>>m[ii][jj*3]>>m[ii][jj*3+1]>>m[ii][jj*3+2];
79     }
80     archivo2.getline(buffer,100);
81 }
82 }
83
84 int main()
85 {
86     int width=225, height=225;
87     float m[225][225*3];
88     leer_data("image.dat",m);
89     GrayScale(m,width,height);
90
91     return EXIT_SUCCESS;
92 }
```

Código 3: Grayscale




Figura 1: Lena original

### 3.2. Blur

```
1 #include <stdio.h>
2 #include <fstream>
3 #include <iostream>
```






Figura 2: Lena a escala de grises

```
4 #define BLUR_SIZE 3
5 using namespace std;
6
7 --global--
8 void blurKernel(float * in, float * out, int w, int h)
9 {
10     int Col = blockIdx.x * blockDim.x + threadIdx.x;
11     int Row = blockIdx.y * blockDim.y + threadIdx.y;
12     if (Col < w && Row < h)
13     {
14         int pixVal = 0;
15         int pixels = 0;
16         // Get the average of the surrounding 2xBLUR_SIZE x 2xBLUR_SIZE box
17         for(int blurRow = -BLUR_SIZE; blurRow < BLUR_SIZE+1; ++blurRow)
18         {
19             for(int blurCol = -BLUR_SIZE; blurCol < BLUR_SIZE+1; ++blurCol)
20             {
21                 int curRow = Row + blurRow;
22                 int curCol = Col + blurCol;
23                 // Verify we have a valid image pixel
24                 if(curRow > -1 && curRow < h && curCol > -1 && curCol < w)
25                 {
26                     pixVal += in[curRow * w + curCol];
27                     pixels++; // Keep track of number of pixels in the
28                             // accumulated total
29                 }
30             }
31         }
32     }
33 }
```

```

30     }
31     // Write our new pixel value out
32     out[Row * w + Col] = (float)(pixVal / pixels);
33 }
34 }
35
36 void save_data(float r[225][225], float g[225][225], float b[225][225])
37 {
38     ofstream archivo("bluur.dat");
39     for (int i = 0; i < 225; ++i)
40     {
41         for (int j = 0; j < 225; ++j)
42         {
43             archivo<<r[i][j]<<" "<<g[i][j]<<" "<<b[i][j]<<" ";
44         }
45         archivo<<endl;
46     }
47 }
48
49 void Blur(float r[225][225], float g[225][225], float b[225][225], int
50 width, int height)
51 {
52     float o_r[225][225];
53     float o_g[225][225];
54     float o_b[225][225];
55
56     int size = width * height;
57     int memSize = size * sizeof(float);
58
59     float *d_A, *d_B;
60
61     cudaMalloc((void **) &d_A, memSize);
62     cudaMalloc((void **) &d_B, memSize);
63
64     cudaMemcpy(d_A, r, memSize, cudaMemcpyHostToDevice);
65     dim3 DimGrid(floor((width-1)/16 + 1), floor((height-1)/16+1), 1);
66     dim3 DimBlock(16, 16, 1);
67     blurKernel<<<DimGrid,DimBlock>>>(d_A, d_B, width, height);
68     cudaMemcpy(o_r, d_B, memSize, cudaMemcpyDeviceToHost);
69
70     cudaMemcpy(d_A, g, memSize, cudaMemcpyHostToDevice);
71
72     blurKernel<<<DimGrid,DimBlock>>>(d_A, d_B, width, height);
73     cudaMemcpy(o_g, d_B, memSize, cudaMemcpyDeviceToHost);
74
75     cudaMemcpy(d_A, b, memSize, cudaMemcpyHostToDevice);
76
77

```

```
78 blurKernel<<<<DimGrid,DimBlock>>>>(d_A, d_B, width, height);
79 cudaMemcpy(o_b, d_B, memSize, cudaMemcpyDeviceToHost);
80
81 cudaFree(d_A);
82 cudaFree(d_B);
83 save_data(o_r, o_g, o_b);
84 }
85
86 void leer_data(const char *file, float r[225][225], float g[225][225],
87               float b[225][225])
88 {
89     char buffer[100];
90     ifstream archivo2("lena.dat");
91     for (int ii = 0; ii < 225; ++ii)
92     {
93         for (int jj = 0; jj < 225; ++jj)
94         {
95             archivo2>>r[ii][jj]>>g[ii][jj]>>b[ii][jj];
96         }
97         archivo2.getline(buffer, 100);
98     }
99 }
100 int main()
101 {
102     int width=225, height=225;
103     float r[225][225];
104     float g[225][225];
105     float b[225][225];
106     leer_data("lena.dat", r, g, b);
107     Blur(r, g, b, width, height);
108     return EXIT_SUCCESS;
109 }
```

Código 4: Blur




Figura 3: Lena original




Figura 4: Lena con BlurSize de 3